

QR Decomposition on ARM920T

By Lucas Antonsen and Ty Ellison

Outline

Introduction to QR Decomposition

Design Challenges

Work Plan

Optimizations

Assembly

Numerical Results

Conclusion

What is QR Decomposition?

$$\begin{array}{c} A \\ \begin{pmatrix} 2.5 & 1.1 & 0.3 & 2.2 \\ 1.9 & 0.4 & 1.8 & 0.1 \\ 0.3 & 0.3 & 0.2 & 1.6 \\ 1.1 & 1. & 0.2 & 1.2 \end{pmatrix} \end{array} = \begin{array}{c} Q \\ \begin{pmatrix} -0.7 & 0.1 & 0.7 & -0.1 \\ -0.6 & -0.6 & -0.6 & 0.1 \\ -0.1 & 0.3 & -0.3 & -0.9 \\ -0.3 & 0.8 & -0.4 & 0.4 \end{pmatrix} \end{array} \begin{array}{c} R \\ \begin{pmatrix} -3.3 & -1.4 & -1.3 & -2.2 \\ 0. & 0.7 & -0.8 & 1.4 \\ 0. & 0. & -1. & 0.4 \\ 0. & 0. & 0. & -1.3 \end{pmatrix} \end{array}$$

Classical Gram-Schmidt Orthogonalization

$$A = [e_1|e_2|\cdots|e_n] \begin{bmatrix} v_1 \cdot e_1 & v_2 \cdot e_1 & \cdots & v_n \cdot e_1 \\ 0 & v_2 \cdot e_2 & \cdots & v_n \cdot e_2 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & v_n \cdot e_n \end{bmatrix} = QR$$

Classical Gram-Schmidt Orthogonalization

$$\mathbf{u}_1 = \mathbf{v}_1,$$

$$\mathbf{u}_2 = \mathbf{v}_2 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_2),$$

$$\mathbf{u}_3 = \mathbf{v}_3 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_3) - \text{proj}_{\mathbf{u}_2}(\mathbf{v}_3),$$

$$\mathbf{u}_4 = \mathbf{v}_4 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_4) - \text{proj}_{\mathbf{u}_2}(\mathbf{v}_4) - \text{proj}_{\mathbf{u}_3}(\mathbf{v}_4),$$

$$\vdots$$

$$\mathbf{u}_k = \mathbf{v}_k - \sum_{j=1}^{k-1} \text{proj}_{\mathbf{u}_j}(\mathbf{v}_k),$$

$$\mathbf{e}_1 = \frac{\mathbf{u}_1}{\|\mathbf{u}_1\|}$$

$$\mathbf{e}_2 = \frac{\mathbf{u}_2}{\|\mathbf{u}_2\|}$$

$$\mathbf{e}_3 = \frac{\mathbf{u}_3}{\|\mathbf{u}_3\|}$$

$$\mathbf{e}_4 = \frac{\mathbf{u}_4}{\|\mathbf{u}_4\|}$$

$$\vdots$$

$$\mathbf{e}_k = \frac{\mathbf{u}_k}{\|\mathbf{u}_k\|}.$$

Modified Gram-Schmidt Orthogonalization

$$\mathbf{u}_k = \mathbf{v}_k - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_k) - \text{proj}_{\mathbf{u}_2}(\mathbf{v}_k) - \cdots - \text{proj}_{\mathbf{u}_{k-1}}(\mathbf{v}_k),$$

$$\mathbf{u}_k^{(1)} = \mathbf{v}_k - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_k),$$

$$\mathbf{u}_k^{(2)} = \mathbf{u}_k^{(1)} - \text{proj}_{\mathbf{u}_2}(\mathbf{u}_k^{(1)}),$$

$$\vdots$$

$$\mathbf{u}_k^{(k-2)} = \mathbf{u}_k^{(k-3)} - \text{proj}_{\mathbf{u}_{k-2}}(\mathbf{u}_k^{(k-3)}),$$

$$\mathbf{u}_k^{(k-1)} = \mathbf{u}_k^{(k-2)} - \text{proj}_{\mathbf{u}_{k-1}}(\mathbf{u}_k^{(k-2)}),$$

$$\mathbf{e}_k = \frac{\mathbf{u}_k^{(k-1)}}{\|\mathbf{u}_k^{(k-1)}\|}$$

CGSO vs. MGSO

Classical Gram-Schmidt orthogonalization

Let $A_j, j = 1, \dots, n$ be linearly independent vectors.

for $j = 1, 2, \dots, n$

$y = A_j$

for $i = 1, 2, \dots, j - 1$

$r_{ij} = q_i^T A_j$

$y = y - r_{ij}q_i$

end

$r_{jj} = \|y\|_2$

$q_j = y/r_{jj}$

end

Modified Gram-Schmidt orthogonalization

Let $A_j, j = 1, \dots, n$ be linearly independent vectors.

for $j = 1, 2, \dots, n$

$y = A_j$

for $i = 1, 2, \dots, j - 1$

$r_{ij} = q_i^T y$

$y = y - r_{ij}q_i$

end

$r_{jj} = \|y\|_2$

$q_j = y/r_{jj}$

end

Images reproduced from the third edition of “Numerical Analysis,” by Timothy Sauer.

Stability GSO vs MGSO




Figure reproduced from [Gram-Schmidt vs. Modified Gram-Schmidt](#) | Laurent Hoeltgen

Properties of QR Decomposition

$$Q^T Q = I$$

$$R = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 0 & 5 & 6 & 7 \\ 0 & 0 & 8 & 9 \\ 0 & 0 & 0 & 10 \end{bmatrix}$$

What is QR Decomposition used for?

- Linear least squares
- Basis for QR Eigenvalue algorithm
- Absolute value of a determinant of a square matrix
- Rank of matrix A

Design Challenge

- Limitations of sequential processors
- Efficient vector and matrix operations
- Fractional arithmetic without a dedicated floating point unit
- Mathematical routines without library support

Code Optimization Plan

- Loop unrolling
- Cache-oblivious matrix transpose
- Macros
- Operator strength reduction
- Inline functions
- Fixed Point Arithmetic
- Reduce branch and load/store instructions
- Dedicated square root instruction

Base Code

- 24 functions
- 20 / 24 functions are 1 or more for loops
- Most functions perform some kind of simple vector operation e.g. element wise addition and subtraction, dot product, square root, and absolute value
- More complex functions require the use of multiple vector routines such as the Euclidean Norm which performs a dot product and square root

Base Code

```
void QR(SIZE_T rows, SIZE_T cols, DATA_T At[cols][rows], NUM_T Q[rows][rows], NUM_T R[rows][cols]) {
    assert(rows >= cols);
    NUM_T y[rows], q[rows];
    NUM_T y_norm;
    SIZE_T i, j;

    for (j=0; j<cols; ++j) {
        vec_copy(rows, At[j], y);

        for (i=0; i<j; ++i) {
            numt_copy_col(rows, cols, i, Q, q);
            R[i][j] = vec_dot(rows, q, y);
            vec_mulc(rows, q, R[i][j]);
            vec_sub(rows, q, y);
        }
        y_norm = l2_norm(rows, y);
        R[j][j] = y_norm;
        vec_divc(rows, y, y_norm);
        numt_set_col(rows, rows, j, y, Q);
    }
}
```

Base Code

```
NUM_T sqr_rt(NUM_T x, NUM_T eps, NUM_T tol, size_t max_iter) {
    assert((int)x >= 0);
    NUM_T x0 = (NUM_T) closest_perfect_square(x, MAX_ITER);
    NUM_T xn, err;
    NUM_T f_n, f_prime;
    size_t i;

    for (i=0; i<max_iter; ++i) {
        f_n = (x0 * x0) - x;
        f_prime = 2. * x0;

        if (abs_val(f_prime) < eps) {
            xn = 0.;
            break;
        }

        xn = x0 - (f_n / f_prime);
        err = abs_val(xn - x0);
        if (err <= tol) {
            break;
        }
        x0 = xn;
    }
    return xn;
}
```

For Loop Example

```
void numt_copy_col(SIZE_T rows, SIZE_T cols, SIZE_T target_col, NUM_T src[rows][cols], NUM_T dest[rows]) {  
    SIZE_T i;  
    for (i=0; i<rows; ++i) {  
        dest[i] = src[i][target_col];  
    }  
}
```



```

NUM_T sqr_rt(NUM_T x, NUM_T eps, NUM_T tol, size_t max_iter) {
    assert((int)x >= 0);
    NUM_T x0 = (NUM_T) closest_perfect_square(x, MAX_ITER);
    NUM_T xn, err;
    NUM_T f_n, f_prime;
    size_t i;

    for (i=0; i<max_iter; ++i) {
        f_n = (x0 * x0) - x;
        f_prime = 2. * x0;

        if (abs_val(f_prime) < eps) {
            xn = 0.;
            break;
        }

        xn = x0 - (f_n / f_prime);
        err = abs_val(xn - x0);
        if (err <= tol) {
            break;
        }
        x0 = xn;
    }
    return xn;
}

```

```

int closest_perfect_square(DATA_T x, size_t max_iter) {
    int sq = 0, xn = 1;
    size_t i;

    for (i=0; i<max_iter; ++i) {
        sq = xn * xn;
        if (sq > (int)x) {
            break;
        }
        xn += 1;
    }
    return xn;
}

```

Optimization 1: Loop Unrolling

```
void vec_sub(SIZE_T size, NUM_T v1[size], NUM_T v2[size]) {  
    SIZE_T i;  
  
    for (i=0; i<size; ++i) {  
        v2[i] -= v1[i];  
    }  
}
```

```
//loop unrolling done  
void vec_sub(SIZE_T size, NUM_T v1[size], NUM_T v2[size]) {  
    SIZE_T i;  
  
    for(i=0; i<size; i+=2){  
        v2[i] -= v1[i];  
  
        if(i+1 != size){  
            v2[i+1] -= v1[i+1];  
        }  
    }  
}
```

```
//loop unrolling done
```

```
void numt_set_col(SIZE_T rows, SIZE_T cols, SIZE_T target_col, NUM_T v[rows], NUM_T A[rows][cols]) {  
    SIZE_T i;  
  
    for (i=0; i<rows; i+=2) {  
        A[i][target_col] = (NUM_T) v[i];  
  
        if(i+1 != rows){  
            A[i+1][target_col] = (NUM_T) v[i+1];  
        }  
    }  
}
```

```
//loop unrolling done
```

```
void mat_set_col(SIZE_T rows, SIZE_T cols, SIZE_T target_col, DATA_T v[rows], DATA_T A[rows][cols]) {  
    SIZE_T i;  
  
    for(i=0; i<rows; i+=2){  
        A[i][target_col] = v[i];  
  
        if(i+1 != rows){  
            A[i+1][target_col] = v[i+1];  
        }  
    }  
}
```

Optimization 2: Cache Oblivious Matrix Transpose

```
void transpose_m(SIZE_T rows, SIZE_T cols, DATA_T A[rows][cols], DATA_T B[cols][rows]) {  
    SIZE_T i, j;  
    for (i=0; i<cols; ++i) {  
        for (j=0; j<rows; ++j) {  
            B[i][j] = A[j][i];  
        }  
    }  
}
```

```

//cache friendly matrix transpose
void transpose_m(SIZE_T rows, SIZE_T cols, DATA_T A[rows][cols], DATA_T B[cols][rows], SIZE_T start_i, SIZE_T end_i, SIZE_T start_j, SIZE_T end_j){

    SIZE_T l_i = end_i - start_i;
    SIZE_T l_j = end_j - start_j;
    SIZE_T i, j;

    if (l_i <= 2 && l_j <= 2) {
        for (i = start_i; i < end_i; i++) {
            for (j = start_j; j < end_j; j++) {
                B[j][i] = A[i][j];
            }
        }
    } else if (l_i >= l_j) {
        transpose_m(rows, cols, A, B, start_i, start_i + (l_i / 2), start_j, end_j);
        transpose_m(rows, cols, A, B, start_i + (l_i / 2), end_i, start_j, end_j);
    } else {
        transpose_m(rows, cols, A, B, start_i, end_i, start_j, start_j + (l_j / 2));
        transpose_m(rows, cols, A, B, start_i, end_i, start_j + (l_j / 2), end_j);
    }
}

```

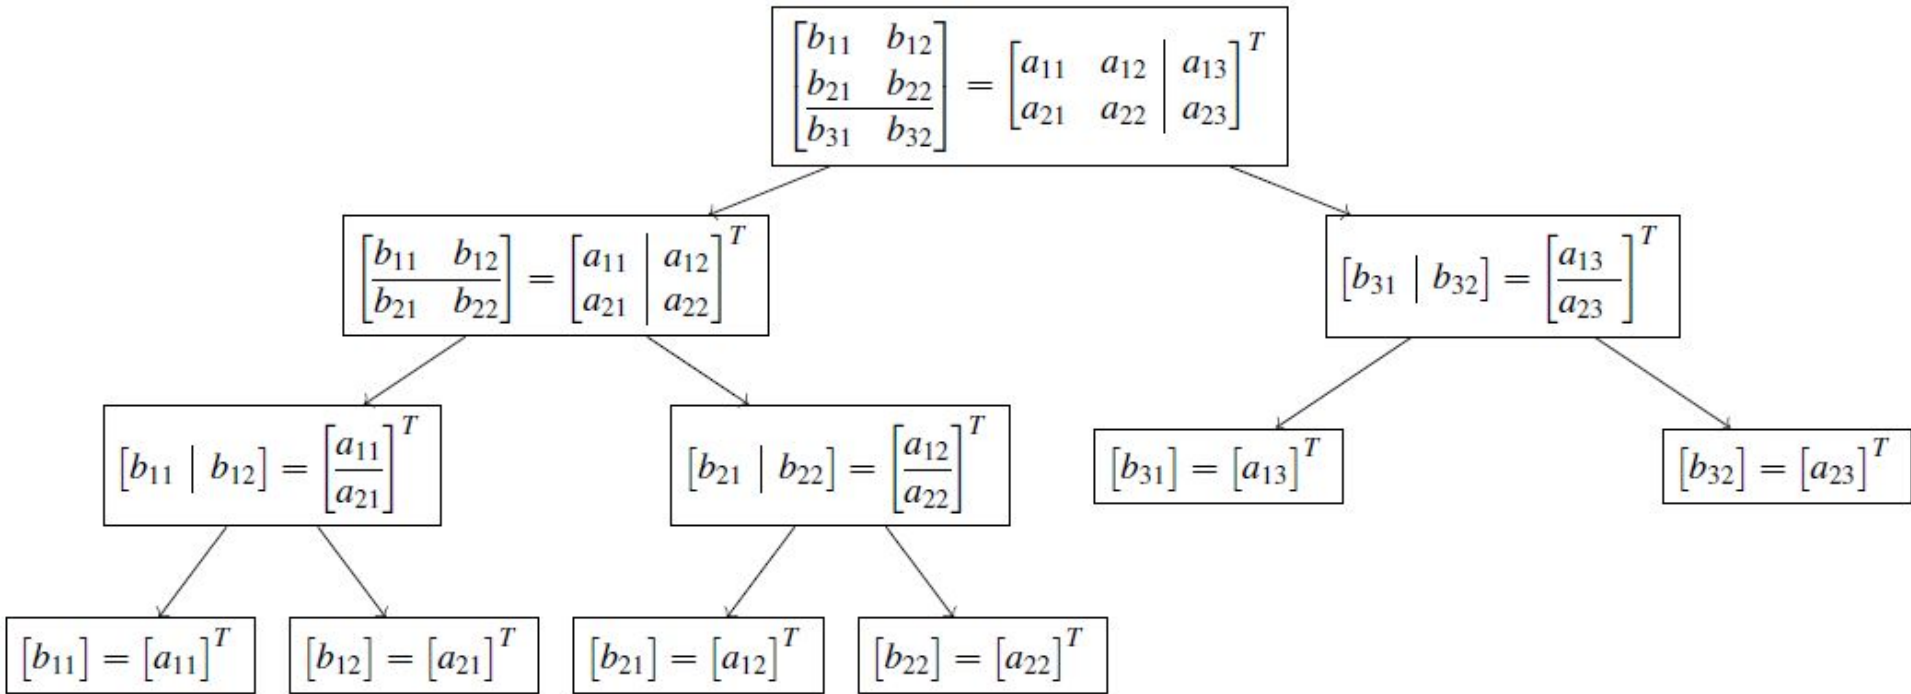



Image from SENG 475 Lecture Slides by Michael Adams

Optimization 3 - Fixed Point Arithmetic



Fixed Point Arithmetic - String to Fixed Width Integer

```
UFX_T str_to_fx(char *s, char *delim, FX_SIZE_T scale) {
    FX_SIZE_T sign = 0;
    FX_T num = 0;
    FX_SIZE_T digits = 0;
    UFX_T threshold = 0;

    char *p;
    char *tok = strtok(s, delim);
    num = (FX_T) strtol(tok, &p, 10);
    if (num < 0) {
        sign = 1;
        num -= 1;
        num = ~num;
    }
    num = ((UFX_T) num << scale);

    tok = strtok(NULL, delim);
    if (tok != NULL) {
        digits = strlen(tok, FX_MAX_DEC_CHARS);
        threshold = get_threshold(tok, digits);
        num |= fract_dec_to_bin((UFX_T) strtoul(tok, &p, 10), threshold);
    }
    return (sign) ? FX_SIGN | num : num;
}
```

Fixed Point Arithmetic - Fixed Width Integer to String

```
void bin_fx_to_str(char *s, UFX_T x) {
    FX_SIZE_T max_chars = FX_MAX_BIN_CHARS - 1;
    UFX_T stack = 0;
    FX_SIZE_T i = 0;

    for (; i < FX_SIZE; ++i) {
        stack <<= 1;
        stack |= (x & 1);
        x >>= 1;
    }

    i = 0;

    if (stack & 1) {
        s[0] = '-';
        i = 1;
        ++max_chars;
    }

    s[FX_WHOLE_BITS + i + 1] = '.';

    for (; i < max_chars; ++i) {
        if ('.' == s[i]) continue;
        s[i] = '0' + (stack & 1);
        stack >>= 1;
    }

    s[FX_MAX_BIN_CHARS] = '\0';
}
```

Fixed Point Arithmetic - Decimal to Binary fractions

```
inline UFX_T fract_dec_to_bin(UFX_T x, UFX_T threshold) {
    UFX_T fract_bin = 0;
    FX_SIZE_T i = 0;

    for (; i < FX_FRACT_BITS; ++i) {
        fract_bin <= 1;
        if ((x <= 1) > threshold) {
            fract_bin |= 1;
            x -= threshold;
        }
    }
    return fract_bin;
}
```

Assembly Comparison - vec_sub Base Version

```
vec_sub:
    @ Function supports interworking.
    @ args = 0, pretend = 0, frame = 0
    @ frame_needed = 0, uses_anonymous_args = 0
    cmp     r0, #0
    stmfd   sp!, {r4, r5, r6, r7, r8, lr}
    mov     r7, r1
    beq     .L31
    sub     r3, r0, #1
    mov     r3, r3, asl #16
    mov     r3, r3, lsr #13
    mov     r5, r2
    add     r6, r3, #8
    mov     r4, #0
```

```
.L30:
    add     r1, r7, r4
    ldmia   r1, {r2-r3}
    ldmia   r5, {r0-r1}
    bl      __aeabi_dsub
    add     r4, r4, #8
    cmp     r4, r6
    stmia   r5!, {r0-r1}
    bne     .L30

.L31:
    ldmfdd  sp!, {r4, r5, r6, r7, r8, lr}
    bx      lr
    .size   vec_sub, .-vec_sub
    .align  2
    .global numt_vec_copy
    .type   numt_vec_copy, %function
```

Assembly Comparison - vec_sub Loop Unrolled Version

```
vec_sub:
    @ Function supports interworking.
    @ args = 0, pretend = 0, frame = 0
    @ frame_needed = 0, uses_anonymous_args = 0
    stmfd    sp!, {r4, r5, r6, r7, r8, r9, s1, lr}
    subs     s1, r0, #0
    mov      r8, r1
    mov      r7, r2
    beq      .L36
    mov      r5, #0


.L35:
    mov      r4, r5, asl #3
    add      r1, r8, r4
    add      r4, r7, r4
    ldmbia   r1, {r2-r3}
    ldmbia   r4, {r0-r1}
    bl       __aeabi_dsub
    add      r2, r5, #1
    mov      r3, r2, asl #3
    cmp      r2, s1

    add      ip, r8, r3
    add      r6, r7, r3
    stmia     r4, {r0-r1}
    beq      .L34
    ldmbia   ip, {r2-r3}
    ldmbia   r6, {r0-r1}
    bl       __aeabi_dsub
    stmia     r6, {r0-r1}

.L34:
    add      r3, r5, #2
    mov      r3, r3, asl #16
    mov      r5, r3, lsr #16
    cmp      s1, r5
    bhi      .L35

.L36:
    ldmfd     sp!, {r4, r5, r6, r7, r8, r9, s1, lr}
    bx        lr
    .size     vec_sub, .-vec_sub
    .align    2
    .global   numt_vec_copy
    .type     numt_vec_copy, %function
```

Dedicated square root instruction



Compile and Run Commands

```
arm-linux-gcc -static -O2 qr.c -o qr.exe
```

```
./qr.exe m20x20.txt out.txt 20 20
```

Numerical Results - Run 100 Times for Average

Base version runtime: 0.739600 seconds

Loop unrolling added: 0.736900 seconds

For loop optimization added: 0.736600 seconds

Cache-oblivious matrix transpose added: 0.736400 seconds

Macros added: 0.736600 seconds

Operator strength reduction added: 0.738400 seconds

Unnecessary loop branches removed: 0.745700 seconds

Matrix transpose run once with 20x20 matrix

I refs: 407,515
I1 misses: 1,013
LLi misses: 1,007
I1 miss rate: 0.25%
LLi miss rate: 0.25%

D refs: 141,806 (89,773 rd + 52,033 wr)
D1 misses: 3,318 (2,518 rd + 800 wr)
LLd misses: 2,628 (1,995 rd + 633 wr)
D1 miss rate: 2.3% (2.8% + 1.5%)
LLd miss rate: 1.9% (2.2% + 1.2%)

LL refs: 4,331 (3,531 rd + 800 wr)
LL misses: 3,635 (3,002 rd + 633 wr)
LL miss rate: 0.7% (0.6% + 1.2%)

Branches: 79,195 (78,038 cond + 1,157 ind)
Mispredicts: 5,700 (5,603 cond + 97 ind)
Mispred rate: 7.2% (7.2% + 8.4%)

I refs: 430,594
I1 misses: 1,017
LLi misses: 1,009
I1 miss rate: 0.24%
LLi miss rate: 0.23%

D refs: 156,878 (98,684 rd + 58,194 wr)
D1 misses: 3,324 (2,517 rd + 807 wr)
LLd misses: 2,630 (1,994 rd + 636 wr)
D1 miss rate: 2.1% (2.6% + 1.4%)
LLd miss rate: 1.7% (2.0% + 1.1%)

LL refs: 4,341 (3,534 rd + 807 wr)
LL misses: 3,639 (3,003 rd + 636 wr)
LL miss rate: 0.6% (0.6% + 1.1%)

Branches: 80,416 (79,259 cond + 1,157 ind)
Mispredicts: 5,878 (5,781 cond + 97 ind)
Mispred rate: 7.3% (7.3% + 8.4%)

Conclusion

- QR decomposition is challenging to optimize when there are many function calls or loops required
- Cache-oblivious matrix transpose lowers cache misses but does not lower runtime
- -O2 provides significant optimization
- FPA is challenging to implement

Future Work

- MLA instruction to parallelize vector and matrix routines
- Implement more robust rounding schemes for decimal to fixed point conversion to reduce round-off error
- Integrate existing QR decomposition code with custom Fixed Point library
- More robust testing procedures which measure round-off and truncation error, and runtime for matrices of arbitrary size and conditioning
- Research other QR decomposition algorithms